

Project 2 (Fall 2025)

AI-Assisted Weight Initialization

Mohamed Ali Kabiri

*Under the supervision of Dr. El Mostafa Kalmoun
Al Akhawayn University in Ifrane*

December 11, 2025

Abstract

This project tested if an AI's advice about neural network initialization is correct. We tested three methods: orthogonal, Gaussian, and uniform initialization. We found that orthogonal initialization works perfectly, Gaussian with $\sigma = 1/\sqrt{d}$ works well, but Gaussian with $\sigma = 1/d$ fails completely. The AI was mostly right but made one big mistake.

1 Introduction

When training neural networks, how you start (initialization) matters a lot. If done wrong, signals can vanish or explode. An AI (Claude) gave us advice on three methods. We tested if that advice was mathematically correct.

2 What We Tested

We tested a 3-layer network: $y = W_3W_2W_1x$ with $d = 50, 100, 200$.

Three initialization methods:

1. **Orthogonal:** W from QR decomposition of random matrix
2. **Gaussian:** $W_{ij} \sim N(0, \sigma^2)$ with different σ
3. **Uniform:** $W_{ij} \sim U(-a, a)$ with $a = \sqrt{3/d}$

We asked Claude: "*Propose three initialization strategies for a 3-layer linear network $y = W_3W_2W_1x$ with $d = 100$: (i) orthogonal, (ii) scaled Gaussian, (iii) random uniform. For each, state a theoretical reason why it should preserve or control the output variance and gradient norms across layers.*"

3 What the AI Said vs. What We Found

3.1 Orthogonal Initialization

AI Claim: "Orthogonal matrices preserve norms exactly, all singular values = 1, prevents vanishing gradients."

Our Test: - Condition number: 1.000 - Singular values: all 1.000 - Gradient ratios: 1.000, 1.000, 1.000 - Variance preservation: 100%

Verdict: Correct!

3.2 Gaussian Initialization

AI Claim 1: " $\sigma = 1/d$ preserves variance."

Our Test ($d = 100$): - Variance preservation: 0.0% (completely vanished!) - Gradient ratios: 0.1006, 0.0088, 0.0009 (vanishing) - Condition number: 710.5 (bad)

Verdict: Wrong!

AI Claim 2: " $\sigma = 1/\sqrt{d}$ preserves variance."

Our Test ($d = 100$): - Variance preservation: 100.4% - Gradient ratios: 0.9533, 0.8885, 0.8761 (good) - Condition number: 158.1 (okay)

Verdict: Correct!

3.3 Uniform Initialization

AI Claim: " $a = \sqrt{3/d}$ matches Gaussian variance."

Our Test ($d = 100$): - Expected variance: 0.01 - Actual variance: 0.010007 - Variance preservation: 97.7% - Gradient ratios: 1.0308, 1.0610, 1.0668 (good)

Verdict: Correct!

4 Data and Results

4.1 Table 1: Matrix Properties ($d = 100$)

Method	Cond. #	Min σ	Max σ	Norm
Orthogonal	1.000	1.000	1.000	1.000
Gaussian $\sigma = 1/d$	710.5	0.000	0.193	0.193
Gaussian $\sigma = 1/\sqrt{d}$	158.1	0.012	1.899	1.899
Uniform	615.0	0.003	2.010	2.010




Figure 1: Singular values for different initializations (d=100)

4.2 Table 2: Variance Propagation ($d = 100$)

Method	Var(z1)	Var(z2)	Var(y)	Preserve
Orthogonal	0.9993	0.9993	0.9993	100.0%
Gaussian $\sigma = 1/d$	0.0101	0.0001	0.0000	0.0%
Gaussian $\sigma = 1/\sqrt{d}$	1.0053	1.0153	1.0094	100.4%
Uniform	0.9864	0.9787	0.9635	97.7%




Figure 2: Variance propagation across layers (d=100)

4.3 Table 3: Gradient Norms ($d = 100$)

Method	$\ g_2\ /\ g\ $	$\ g_1\ /\ g\ $	$\ g_0\ /\ g\ $
Orthogonal	1.0000	1.0000	1.0000
Gaussian $\sigma = 1/d$	0.1006	0.0088	0.0009
Gaussian $\sigma = 1/\sqrt{d}$	0.9533	0.8885	0.8761
Uniform	1.0308	1.0610	1.0668




Figure 3: Gradient norm ratios during backpropagation ($d=100$)

5 Code Summary

We wrote MATLAB code to:

- Generate orthogonal matrices using QR decomposition
- Generate Gaussian matrices with different σ
- Generate uniform matrices with $a = \sqrt{3/d}$
- Test variance propagation with 2000 random inputs
- Test gradient flow with random gradients
- Compute condition numbers and singular values

All code is included in the ZIP file.

6 Conclusion

6.1 What We Learned

- Orthogonal initialization works perfectly - preserves everything exactly

- Gaussian with $\sigma = 1/\sqrt{d}$ works well - good for practical use
- Gaussian with $\sigma = 1/d$ fails completely - don't use this!
- Uniform works like Gaussian $\sigma = 1/\sqrt{d}$ - also good

6.2 Was the AI Correct?

- **Correct:** 7 claims (orthogonal properties, $\sigma = 1/\sqrt{d}$, uniform)
- **Mostly correct:** 3 claims (needed more detail)
- **Wrong:** 2 claims ($\sigma = 1/d$ is bad!)

6.3 Final Advice

1. It's clear that we need to use orthogonal if you want guaranteed stability
2. And use Gaussian $\sigma = 1/\sqrt{d}$ or Uniform for most cases
3. Also, never use Gaussian $\sigma = 1/d$

Appendix: Mathematical Theorems Used

1. Properties of Orthogonal Matrices

If Q is orthogonal ($Q^T Q = I$), then:

- Norm preservation: $\|Qx\| = \|x\|$ for any vector x
- All singular values = 1
- Condition number $\kappa(Q) = 1$
- $Q^{-1} = Q^T$

2. Variance of Linear Transforms

For $z = Wx$ where $x \sim \mathcal{N}(0, I)$ and $W_{ij} \sim \mathcal{N}(0, \sigma^2)$:

$$E[\|z\|^2] = \sum_{j=1}^d E[z_j^2] = d\sigma^2 E[\|x\|^2]$$

To preserve variance: $d\sigma^2 = 1 \Rightarrow \sigma = 1/\sqrt{d}$.

3. Product Matrix Condition Number

For $P = W_3W_2W_1$:

$$\kappa(P) \leq \kappa(W_3) \cdot \kappa(W_2) \cdot \kappa(W_1)$$

This shows why orthogonal initialization ($\kappa = 1$) remains stable.

Appendix A: MATLAB Code Files

- `orthogonal_init.m` - makes orthogonal matrices
- `gaussian_init.m` - makes Gaussian matrices
- `uniform_init.m` - makes uniform matrices
- `main.m` - runs all tests
- `make_plot.m` - make the graphs

Appendix B: AI Chat Transcript

Note: Dear Dr.Kalmoun, the complete Claude conversation is included as a PDF file in the submission. Below is a summary:

Claude's Main Recommendations:

1. Orthogonal initialization via QR decomposition
2. Gaussian with $\sigma^2 = 1/d$ or $\sigma^2 = 1/\sqrt{d}$
3. Uniform with $a = \sqrt{3/d}$ to match Gaussian variance

Our Critique:

- Correct about orthogonal initialization
- Correct about $\sigma = 1/\sqrt{d}$
- Wrong about $\sigma = 1/d$ (causes vanishing)
- Correct about uniform matching Gaussian

Appendix C: How to Reproduce Results

- MATLAB version: Any recent version (R2024a used)
- Random seed: 12345 (set in code for reproducibility)
- Run `main.m` to generate all results
- Run `part_c_experiments.m` for Part C tests